

MScFE_Capstone_Workbook_M6

October 23, 2024

1 Market Timing Strategies using Kalman Filter

WorldQuant University, Msc Financial Engineering

author: rakeshsharma.pr@gmail.com

2 Introduction

In this section, I will explore the implementation of a Kalman filter in Python using the library Pykalman to forecast stock prices. The 2 stocks we are targeting are Ethereum and Bitcoin. As discussed in the literature review, research shows that Kalman filters have better outcomes for time series forecasting and can be used with various pricing algorithms, such as Heston, Bates, etc.

```
[1]: # installation of libraries
```

```
!pip install pykalman
!pip install plotly
!pip install statsmodels
```

Collecting pykalman

Downloading pykalman-0.9.7-py2.py3-none-any.whl.metadata (5.5 kB)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from pykalman) (1.26.4)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from pykalman) (1.13.1)

Downloading pykalman-0.9.7-py2.py3-none-any.whl (251 kB)

251.6/251.6 kB

3.7 MB/s eta 0:00:00

Installing collected packages: pykalman

Successfully installed pykalman-0.9.7

Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (5.24.1)

Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly) (9.0.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from plotly) (24.1)

Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (0.14.4)

Requirement already satisfied: numpy<3,>=1.22.3 in

```

/usr/local/lib/python3.10/dist-packages (from statsmodels) (1.26.4)
Requirement already satisfied: scipy!=1.9.2,>=1.8 in
/usr/local/lib/python3.10/dist-packages (from statsmodels) (1.13.1)
Requirement already satisfied: pandas!=2.1.0,>=1.4 in
/usr/local/lib/python3.10/dist-packages (from statsmodels) (2.2.2)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-
packages (from statsmodels) (0.5.6)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.10/dist-packages (from statsmodels) (24.1)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-
packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.2)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages
(from patsy>=0.5.6->statsmodels) (1.16.0)

```

```

[2]: # importing of libraries

import pandas as pd
import yfinance as yf
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.graph_objects as go
import statsmodels.api as sm

```

3 Exploratory Data Analysis

```

[3]: # crypto currencies of interest - Bitcoin and Ethereum

```

```

stock_list = ['BTC-USD', 'ETH-USD']

```

```

[4]: # extracting the stock price for the last 2 years
data = yf.download(stock_list, start='2022-09-01', end='2024-08-31')
data.head()

```

```

[*****100%*****] 2 of 2 completed

```

```

[4]: Price                Adj Close                Close \
      Ticker                BTC-USD                ETH-USD                BTC-USD
Date
2022-09-01 00:00:00+00:00  20127.140625  1586.176758  20127.140625
2022-09-02 00:00:00+00:00  19969.771484  1577.220459  19969.771484
2022-09-03 00:00:00+00:00  19832.087891  1556.872681  19832.087891
2022-09-04 00:00:00+00:00  19986.712891  1577.641602  19986.712891

```

```
2022-09-05 00:00:00+00:00 19812.371094 1617.183228 19812.371094
```

```
Price
Ticker          ETH-USD          High          \
Date
2022-09-01 00:00:00+00:00 1586.176758 20198.390625 1593.082764
2022-09-02 00:00:00+00:00 1577.220459 20401.568359 1643.183228
2022-09-03 00:00:00+00:00 1556.872681 20037.009766 1579.454346
2022-09-04 00:00:00+00:00 1577.641602 19999.689453 1578.009277
2022-09-05 00:00:00+00:00 1617.183228 20031.160156 1621.661377
```

```
Price
Ticker          Low          Open          \
Date          BTC-USD          ETH-USD          BTC-USD
2022-09-01 00:00:00+00:00 19653.968750 1520.188354 20050.498047
2022-09-02 00:00:00+00:00 19814.765625 1551.877930 20126.072266
2022-09-03 00:00:00+00:00 19698.355469 1541.672119 19969.718750
2022-09-04 00:00:00+00:00 19636.816406 1543.698853 19832.470703
2022-09-05 00:00:00+00:00 19673.046875 1559.781860 19988.789062
```

```
Price
Ticker          Volume
Date          ETH-USD          BTC-USD          ETH-USD
2022-09-01 00:00:00+00:00 1553.756348 30182031010 16434276817
2022-09-02 00:00:00+00:00 1586.017944 29123998928 17708478709
2022-09-03 00:00:00+00:00 1577.213745 23613051457 9516825994
2022-09-04 00:00:00+00:00 1556.895874 25245861652 8884144998
2022-09-05 00:00:00+00:00 1577.884033 28813460025 13060541168
```

```
[5]: data.describe()
```

```
[5]: Price      Adj Close      Close      High      \
Ticker      BTC-USD      ETH-USD      BTC-USD      ETH-USD      BTC-USD
count      730.000000      730.000000      730.000000      730.000000      730.000000
mean      37518.319863      2154.649833      37518.319863      2154.649833      38162.137949
std       17685.270401      760.858790      17685.270401      760.858790      18083.910703
min       15787.284180      1100.169800      15787.284180      1100.169800      16253.047852
25%       23665.855469      1620.496521      23665.855469      1620.496521      24119.581543
50%       29412.204102      1866.100159      29412.204102      1866.100159      29845.836914
75%       55988.014648      2640.965576      55988.014648      2640.965576      57679.622070
max       73083.500000      4066.445068      73083.500000      4066.445068      73750.070312
```

```
Price
Ticker          Low          Open          \
Date          ETH-USD          BTC-USD          ETH-USD          BTC-USD          ETH-USD
count      730.000000      730.000000      730.000000      730.000000      730.000000
mean      2197.880829      36780.942348      2107.257615      37464.146091      2153.339377
std       781.196683      17220.739292      736.665024      17678.426052      761.057358
```

min	1136.442627	15599.046875	1081.138184	15782.300781	1100.107178
25%	1644.864014	23253.754883	1580.678528	23627.717285	1617.854645
50%	1887.949524	29113.966797	1845.784241	29403.917969	1865.844604
75%	2709.429199	54234.083008	2585.498413	55644.687500	2640.590576
max	4092.284180	71334.093750	3936.627197	73079.375000	4066.690430

	Price	Volume
Ticker	BTC-USD	ETH-USD
count	7.300000e+02	7.300000e+02
mean	2.494924e+10	1.111478e+10
std	1.410581e+10	6.917282e+09
min	5.331173e+09	2.081626e+09
25%	1.473525e+10	6.289494e+09
50%	2.165289e+10	9.412596e+09
75%	3.145988e+10	1.414162e+10
max	1.189925e+11	6.766813e+10

The data shows that open, close, High,low are correlated, we can check this in our plot

```
[6]: data.corr().style.background_gradient(cmap='coolwarm')
```

```
[6]: <pandas.io.formats.style.Styler at 0x78ef7d123550>
```

```
[7]: data.isnull().any()
```

```
[7]: Price      Ticker
Adj Close  BTC-USD  False
           ETH-USD  False
Close      BTC-USD  False
           ETH-USD  False
High       BTC-USD  False
           ETH-USD  False
Low        BTC-USD  False
           ETH-USD  False
Open       BTC-USD  False
           ETH-USD  False
Volume     BTC-USD  False
           ETH-USD  False
dtype: bool
```

The dataset do not contain any null values

```
[8]: closing_data = data['Adj Close']['BTC-USD']
      closing_data.head()
```

```
[8]: Date
2022-09-01 00:00:00+00:00    20127.140625
2022-09-02 00:00:00+00:00    19969.771484
```

```

2022-09-03 00:00:00+00:00    19832.087891
2022-09-04 00:00:00+00:00    19986.712891
2022-09-05 00:00:00+00:00    19812.371094
Name: BTC-USD, dtype: float64

```

```

[9]: closing_data = data['Adj Close']['ETH-USD']
     closing_data.head()

```

```


[9]: Date
2022-09-01 00:00:00+00:00    1586.176758
2022-09-02 00:00:00+00:00    1577.220459
2022-09-03 00:00:00+00:00    1556.872681
2022-09-04 00:00:00+00:00    1577.641602
2022-09-05 00:00:00+00:00    1617.183228
Name: ETH-USD, dtype: float64

```

```

[10]: fig, axs = plt.subplots(1,2,figsize=(16, 5),gridspec_kw ={'hspace': 0.2,
    ↪ 'wspace': 0.1})
     axs[0].plot(data['Adj Close']['BTC-USD'])
     axs[0].set_title('BTC')
     axs[1].plot(data['Adj Close']['ETH-USD'])
     axs[1].set_title('ETH')
     plt.show()

```




Lets Visualize the Volume for trend, noise and seasonality


```

[74]: fig, axs = plt.subplots(1,2,figsize=(16, 5),gridspec_kw ={'hspace': 0.2,
    ↪ 'wspace': 0.1})
     axs[0].plot(data['Volume']['BTC-USD'])
     axs[0].set_title('BTC')
     axs[1].plot(data['Volume']['ETH-USD'])
     axs[1].set_title('ETH')
     plt.show()


```



```
[81]: fig, axs =plt.subplots(1,2,figsize=(16, 5),gridspec_kw ={'hspace': 0.2,␣
      ↪ 'wspace': 0.1})
      axs[0].plot(data['Volume']['BTC-USD'].diff())
      axs[0].set_title('BTC')
      axs[1].plot(data['Volume']['ETH-USD'].diff())
      axs[1].set_title('ETH')
      plt.show()
```




```
[82]: fig, axs =plt.subplots(1,2,figsize=(16, 5),gridspec_kw ={'hspace': 0.2,␣
      ↪ 'wspace': 0.1})
      axs[0].plot(data['Volume']['BTC-USD'].diff().diff())
      axs[0].set_title('BTC')
      axs[1].plot(data['Volume']['ETH-USD'].diff().diff())
      axs[1].set_title('ETH')
      plt.show()
```



```
[79]: fig, axs = plt.subplots(1,2,figsize=(16,5),gridspec_kw ={'hspace': 0.2,
↳ 'wspace': 0.1})
btc = data['Adj Close']['BTC-USD'].pct_change().dropna(axis=0)
axs[0].plot(btc)
axs[0].set_title('BTC')
axs[0].set_ylim([-0.5,0.5])

eth = data['Adj Close']['ETH-USD'].pct_change().dropna(axis=0)
axs[1].plot(eth)
axs[1].set_title('ETH')
axs[1].set_ylim([-0.5,0.5])
```

[79]: (-0.5, 0.5)



OHLC plots - Price movements A candle stick plot helps us understand the spread between the prices of the crypto currency from the opening and closing price and identity the periods where they were continuously on the rise.

```
[12]: fig = go.Figure(data=[go.Candlestick(x=data.index,
open=data['Open']['BTC-USD'], high=data['High']['BTC-USD'],
```

```

        low=data['Low']['BTC-USD'], close=data['Close']['BTC-USD'])
    ])
fig.update_layout(xaxis_rangeslider_visible=True)
fig.show()

```

```

[13]: fig = go.Figure(data=[go.Candlestick(x=data.index,
        open=data['Open']['ETH-USD'], high=data['High']['ETH-USD'],
        low=data['Low']['ETH-USD'], close=data['Close']['ETH-USD'])
    ])
fig.update_layout(xaxis_rangeslider_visible=True)
fig.show()

```

Understanding the Price shocks

We can see how the rolling mean for various intervals looks like for our crypto data, we find that with increasing window size our data gets smoother and allows us to see the trend over which can be used for generating feature engineering. The shorter the window closer it is to the actual stock price and considers the price shocks.


```

[14]: btc_closing_data = data['Adj Close']['BTC-USD'].loc['2024-05-01':'2024-09-01']
rolling_3d = btc_closing_data.rolling(window=3).mean()
rolling_7d = btc_closing_data.rolling(window=7).mean()
rolling_14d = btc_closing_data.rolling(window=14).mean()
rolling_28d = btc_closing_data.rolling(window=28).mean()

plt.figure(figsize=(16, 5))
plt.plot(btc_closing_data.index, btc_closing_data, lw=3, alpha=0.8,
        label='BTC-USD Daily Observation')
plt.plot(btc_closing_data.index, rolling_3d, lw=3, alpha=0.8, label='Rolling
        mean - window 3')
plt.plot(btc_closing_data.index, rolling_7d, lw=3, alpha=0.8, label='Rolling
        mean - window 7')
plt.plot(btc_closing_data.index, rolling_14d, lw=3, alpha=0.8, label='Rolling
        mean - window 14')
plt.plot(btc_closing_data.index, rolling_28d, lw=3, alpha=0.8, label='Rolling
        mean - window 28')

plt.title('BTC-USD Adj Close Price - Observation Window 2024-05-01 :
        2024-09-01')
plt.tick_params(labelsize=12)
plt.legend(loc='upper left', fontsize=12)
plt.show()


```

```
[15]: eth_closing_data = data['Adj Close']['ETH-USD'].loc['2024-05-01':'2024-09-01']
rolling_3d = eth_closing_data.rolling(window=3).mean()
rolling_7d = eth_closing_data.rolling(window=7).mean()
rolling_14d = eth_closing_data.rolling(window=14).mean()
rolling_28d = eth_closing_data.rolling(window=28).mean()

plt.figure(figsize=(16, 5))
plt.plot(eth_closing_data.index, eth_closing_data, lw=3, alpha=0.8, label='ETH-USD Daily Observation')
plt.plot(eth_closing_data.index, rolling_3d, lw=3, alpha=0.8, label='Rolling mean - window 3')
plt.plot(eth_closing_data.index, rolling_7d, lw=3, alpha=0.8, label='Rolling mean - window 7')
plt.plot(eth_closing_data.index, rolling_14d, lw=3, alpha=0.8, label='Rolling mean - window 14')
plt.plot(eth_closing_data.index, rolling_28d, lw=3, alpha=0.8, label='Rolling mean - window 28')

plt.title('ETH-USD Adj Close Price - Observation Window 2024-05-01 : 2024-09-01')
plt.tick_params(labelsize=12)
plt.legend(loc='upper left', fontsize=12)
plt.show()
```



Ethereum rolling means are lot steeper than BTC, in the further sections we will do the causality tests to see autocorrelation.

Understanding the Volume Shocks


Positive volume shocks can be leveraged to improve returns and it is significant for volume based trading in crypto. Trading strategists can buy stocks on a positive volume jump and sell stocks on volume downsides to make profit over a period of time. A shock is defined as the increase or decrease in the volume comapred to the 12-month average. Compared to the rolling mean for 28 days, daily volume shows significant upward spikes compared to downfalls.

```
[16]: btc_closing_data = data['Volume']['BTC-USD'].loc['2024-05-01':'2024-09-01']
rolling_3d = btc_closing_data.rolling(window=3).mean()
rolling_7d = btc_closing_data.rolling(window=7).mean()
rolling_14d = btc_closing_data.rolling(window=14).mean()
rolling_28d = btc_closing_data.rolling(window=28).mean()

plt.figure(figsize=(16, 5))
plt.plot(btc_closing_data.index, btc_closing_data, lw=3, alpha=0.8,
        label='BTC-USD Daily Observation')
plt.plot(btc_closing_data.index, rolling_3d, lw=3, alpha=0.8, label='Rolling
        mean - window 3')
plt.plot(btc_closing_data.index, rolling_7d, lw=3, alpha=0.8, label='Rolling
        mean - window 7')
plt.plot(btc_closing_data.index, rolling_14d, lw=3, alpha=0.8, label='Rolling
        mean - window 14')
plt.plot(btc_closing_data.index, rolling_28d, lw=3, alpha=0.8, label='Rolling
        mean - window 28')

plt.title('BTC-USD Volume Shocks - Observation Window - 2024-05-01 :
        2024-09-01')
plt.tick_params(labelsize=12)
plt.legend(loc='upper left', fontsize=12)
```


```
plt.show()
```



```
[17]: btc_closing_data = data['Volume']['ETH-USD'].loc['2024-05-01':'2024-09-01']
rolling_3d = btc_closing_data.rolling(window=3).mean()
rolling_7d = btc_closing_data.rolling(window=7).mean()
rolling_14d = btc_closing_data.rolling(window=14).mean()
rolling_28d = btc_closing_data.rolling(window=28).mean()

plt.figure(figsize=(16, 5))
plt.plot(btc_closing_data.index, btc_closing_data, lw=3, alpha=0.8,
        label='BTC-USD Daily Observation')
plt.plot(btc_closing_data.index, rolling_3d, lw=3, alpha=0.8, label='Rolling
        mean - window 3')
plt.plot(btc_closing_data.index, rolling_7d, lw=3, alpha=0.8, label='Rolling
        mean - window 7')
plt.plot(btc_closing_data.index, rolling_14d, lw=3, alpha=0.8, label='Rolling
        mean - window 14')
plt.plot(btc_closing_data.index, rolling_28d, lw=3, alpha=0.8, label='Rolling
        mean - window 28')

plt.title('ETH-USD Volume Shocks - Observation Window - 2024-05-01 :
        2024-09-01')
plt.tick_params(labels=12)
plt.legend(loc='upper left', fontsize=12)
plt.show()
```




Check for Autocorrelation and Partial autocorrelation

In the below section we can check for ACF and PACF functions for 'Adjusted Closing Price'

```
[18]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 5))

btc_data = data['Adj Close']['BTC-USD']
sm.graphics.tsa.plot_acf(btc_data.values.squeeze(), lags=30, ax=ax1,
    ↳title="BTC-USD ACF Plot")


eth_data = data['Adj Close']['ETH-USD']
sm.graphics.tsa.plot_acf(eth_data.values.squeeze(), lags=30, ax=ax2,
    ↳title="ETH-USD ACF Plot")
plt.show()
```



```
[19]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 5))

btc_data = data['Adj Close']['BTC-USD']
sm.graphics.tsa.plot_pacf(btc_data.values.squeeze(), lags=30, ax=ax1,
    ↳title="BTC-USD PACF Plot")
```

```
eth_data = data['Adj Close']['ETH-USD']
sm.graphics.tsa.plot_pacf(eth_data.values.squeeze(), lags=30, ax=ax2,
    ↪title="ETH-USD PACF Plot")
plt.show()
```




The plot shows significant autocorrelation, the previous value of the currency determines days closing price. From the PACF plot it is clear that price of certain period of time is correlated than others, say for example price on day 4, 9, 12, 15 etc.

In the next sections we will check for Autocorrelations and Partial Autocorrelations for Volume in both crypto currencies.

```
[20]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 5))

btc_data = data['Volume']['BTC-USD']
sm.graphics.tsa.plot_acf(btc_data.values.squeeze(), lags=30, ax=ax1,
    ↪title="BTC-USD Volume ACF Plot")


eth_data = data['Volume']['ETH-USD']
sm.graphics.tsa.plot_acf(eth_data.values.squeeze(), lags=30, ax=ax2,
    ↪title="ETH-USD Volume ACF Plot")
plt.show()
```



```
[21]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 5))

btc_data = data['Volume']['BTC-USD']
sm.graphics.tsa.plot_pacf(btc_data.values.squeeze(), lags=30, ax=ax1,
    ↳title="BTC-USD Volume PACF Plot")

eth_data = data['Volume']['ETH-USD']
sm.graphics.tsa.plot_pacf(eth_data.values.squeeze(), lags=30, ax=ax2,
    ↳title="ETH-USD Volume PACF Plot")
plt.show()
```



4 Statistical Test for Stationarity

In this section we will use ADF test to establish that dataset is stationary. It is a unit-root test for time series.

```
[72]: from statsmodels.tsa.stattools import adfuller
```

```

print("P-value for AD Fuller test for BTC-USD Adj Close {}".format(adfuller(
    ↳data['Adj Close']['BTC-USD'])[1]))
print("P-value for AD Fuller test for ETH-USD Adj Close {}".format(adfuller(
    ↳data['Adj Close']['ETH-USD'])[1]))
print("P-value for AD Fuller test for BTC-USD Volume {}".format(adfuller(
    ↳data['Volume']['BTC-USD'])[1]))
print("P-value for AD Fuller test for ETH-USD Volume {}".format(adfuller(
    ↳data['Volume']['ETH-USD'])[1]))

```

P-value for AD Fuller test for BTC-USD Adj Close 0.8759619045006749
 P-value for AD Fuller test for ETH-USD Adj Close 0.6087041759596575
 P-value for AD Fuller test for BTC-USD Volume 0.011305878023385845
 P-value for AD Fuller test for ETH-USD Volume 0.02339000739363443

From the AD Fuller test, it can be seen for Adj Close Price the p-value is greater than 0.05 for ETH-USD and BTC-USD, indicating the data series is non-stationary. In the case Volume the p-value is less than 0.05, therefore we reject the null hypothesis and the series is stationary. We can see this with Kalman Forecast below for Volume, they do not show any trend.

5 Exploring Kalman Filter

Getting started with Kalman Filters for the data, we will be using the algorithm's implementation in pykalman. Let us do the Kalman smoothing and filtering methods to remove the trend here. Stock price by nature is very noisy and with an underlying trend. The identification of this trend can be useful in many ways and can also be used as a feature for stock price forecasting using ARIMA, Linear Regression or any other methods.

[22]: *# initializing kalman filters*

```

from pykalman import KalmanFilter
import numpy as np

kalmanFilter = KalmanFilter(transition_matrices = [1],
                           observation_matrices = [1],
                           initial_state_mean = 0,
                           initial_state_covariance = 1,
                           observation_covariance=1,
                           transition_covariance=.0001)

```

[23]: *# using Kalman Filters for BTC-USD*

```

stock_price = data['Adj Close']['BTC-USD']
mean, cov = kalmanFilter.filter(stock_price)
kalman_mean = pd.Series(mean.flatten(), index=stock_price.index)


```

```

rolling_3d = stock_price.rolling(window=3).mean()
rolling_7d = stock_price.rolling(window=7).mean()
rolling_14d = stock_price.rolling(window=14).mean()
rolling_28d = stock_price.rolling(window=28).mean()

plt.figure(figsize=(12,6))
plt.plot(rolling_7d, '-g', lw=1.5)
plt.plot(rolling_14d, 'm', lw=1)
plt.plot(rolling_28d, 'y', lw=1)
plt.plot(kalman_mean, 'r', lw=1)
plt.plot(stock_price, 'b', lw=1)
plt.title('Kalman filter estimate of average for "Adj Close" - BTC-USD')
plt.legend(['7-day Moving Average', '14-day Moving Average', '28day Moving Average', 'Kalman Mean', 'Actual Stock Price'])
plt.xlabel('Day')
plt.ylabel('Adj Close Price');

```



```

[24]: plt.figure(figsize=(15,7))
plt.plot(stock_price, 'b', lw=1.5)
plt.plot(kalman_mean, 'm', lw=1.5)
plt.plot(stock_price - kalman_mean, '-g', lw=1.5)
plt.plot(stock_price - stock_price.mean(), 'r', lw=1.5)
plt.plot(stock_price - rolling_28d, 'y', lw=1.5)
plt.plot(stock_price - rolling_3d, 'k', lw=1.5)
plt.legend(['Stock Price', 'Kalman Mean', 'Diff between Stock Price and Kalman Mean', 'Diff between Stock Price and Mean Price',


```



```

        'Diff between Stock Price and rolling mean of 30 days', 'Diff_
↪between Stock Price and rolling mean of 3 days'])
plt.xlabel('Day')
plt.ylabel('Variance of the Stock Price Prediction with Methods for BTC-USD');

```



```

[25]: kalmanFilter = KalmanFilter(transition_matrices = [1],
    observation_matrices = [1],
    initial_state_mean = 0,
    initial_state_covariance = 1,
    observation_covariance=1,
    transition_covariance=.0001)


stock_price = data['Adj Close']['ETH-USD']
mean, cov = kalmanFilter.filter(stock_price)
kalman_mean = pd.Series(mean.flatten(), index=stock_price.index)

rolling_3d = stock_price.rolling(window=3).mean()
rolling_7d = stock_price.rolling(window=7).mean()
rolling_14d = stock_price.rolling(window=14).mean()
rolling_28d = stock_price.rolling(window=28).mean()

plt.figure(figsize=(15,7))
plt.plot(rolling_7d, '-g', lw=1.5)
plt.plot(rolling_14d, 'm', lw=1)
plt.plot(rolling_28d, 'y', lw=1)
plt.plot(kalman_mean, 'r', lw=1)
plt.plot(stock_price, 'b', lw=1)

```


```
plt.title('Kalman filter estimate of average for "Adj Close" ETH-USD')
plt.legend(['7-day Moving Average', '14-day Moving Average', '28-day Moving Average', 'Kalman Mean', 'Actual Stock Price'])
plt.xlabel('Day')
plt.ylabel('Adj Close Price');
```



From the kalman Filter for Bitcoin it can be seen that it produces much more smooth signal than others and is more generalised than the rolling window prediction. Kalman filter helps in separating the noise from the signals and can be used for prediction

```
[26]: plt.figure(figsize=(15,7))
plt.plot(stock_price , 'b', lw=1.5)
plt.plot(kalman_mean , 'm', lw=1.5)
plt.plot(stock_price - kalman_mean, '-g', lw=1.5)
plt.plot(stock_price - stock_price.mean() , 'r', lw=1.5)
plt.plot(stock_price - rolling_28d , 'y', lw=1.5)
plt.plot(stock_price - rolling_3d , 'k', lw=1.5)

plt.legend(['Stock Price', 'Kalman Mean', 'Diff between Stock Price and Kalman Mean', 'Diff between Stock Price and Mean Price', 'Diff between Stock Price and rolling mean of 30 days', 'Diff between Stock Price and rolling mean of 3 days'])
plt.xlabel('Day')
plt.ylabel('Variance of the Stock Price Prediction with Methods for ETH-USD');
```



Kalman Filters gives a time-evolving mean and shows fluctuations with time than the and difference stays closer to 0 compared to the other steps such as rolling mean and static mean. Such updated mean for each time point could help us manage aggressive trading strategies.

In the below section we will using Kalman Filter update method to do stock price forecasts for each day using the previous data as the prior and will keep on update the price on based on the daily available actual price like a dynamic system.

In the below Sections we will forecast Kalman Filter for **ETH-USD** data for various time intervals. We will be using Kalman Filter update on 'Adj Close' column of the time series.

6 Kalman Filter for 'Adj Close'

```
[27]: kalmanFilter = KalmanFilter(transition_matrices = [1],
                                observation_matrices = [1],
                                initial_state_mean = 0,
                                initial_state_covariance = 1,
                                observation_covariance=1,
                                transition_covariance=.0001)

# Predicting for ETH-USD
stock_price = data['Adj Close']['ETH-USD']

# Kalman Forecast for Next 100 days
mean, cov = kalmanFilter.em(stock_price[:-200], n_iter=10).smooth(stock_price[:
↪-200]) # changed from 10 to 100, after experiment put it back to 10
# mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

next_means = []
```

```

next_covs = []
next_mean = mean[-1]
next_cov = cov[-1]
for i in range(200):
    next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
↪stock_price[-(200 - i * 1)]) # changed from 10 to 100, after experiment put
↪it back to 10
    next_means.append(next_mean[0])
    next_covs.append(next_cov[0])


# replacing the forecasted price for the last 10 days
forecasted_price = pd.DataFrame(data=np.concatenate([stock_price[:-(200 - 10)].values,
↪next_means]),
                                index=stock_price.index)

plt.figure(figsize=(12,6))
plt.plot(stock_price[len(stock_price)-200:], 'b',lw=1.5)
plt.plot(forecasted_price[len(stock_price)-200:], 'r',lw=1.5)
plt.legend(['Stock Price', 'Kalman Filter Forecasted Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price');

```

<ipython-input-27-d5d7869a3599>:21: FutureWarning:

Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`



```
[28]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
      ↪ mean_squared_error, mean_absolute_error
      import math

      _y_stock = stock_price[len(stock_price)-200:] # changed from 10 to 100, after
      ↪ experiment put it back to 10
      kalman_forecast_y = forecasted_price[len(stock_price)-200:] # changed from 10
      ↪ to 100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
      print("Mean absolute percentage error {}".
      ↪ format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
      print("Root Mean Square Error RMSE {}".format(math.
      ↪ sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
      print("Mean absolute error {}".format(mean_absolute_error(_y_stock,
      ↪ kalman_forecast_y)))
```

R square 0.5331947237800019
Mean absolute percentage error 0.06379022509833535
Root Mean Square Error RMSE 262.8734331675915
Mean absolute error 204.2497629668713

```
[29]: kalmanFilter = KalmanFilter(transition_matrices = [1],
      observation_matrices = [1],
      initial_state_mean = 0,
      initial_state_covariance = 1,
      observation_covariance=1,
      transition_covariance=.0001)

      # Predicting for ETH-USD
      stock_price = data['Adj Close']['ETH-USD']

      # Kalman Forecast for Next 100 days
      mean, cov = kalmanFilter.em(stock_price[:-100], n_iter=10).smooth(stock_price[:
      ↪ -100]) # changed from 10 to 100, after experiment put it back to 10
      # mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

      next_means = []
      next_covs = []
      next_mean = mean[-1]
      next_cov = cov[-1]
      for i in range(100):
```

```

next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
↪stock_price[-(100 - i * 1)]) # changed from 10 to 100, after experiment put
↪it back to 10
next_means.append(next_mean[0])
next_covs.append(next_cov[0])

```

<ipython-input-29-968bb0fc92b6>:21: FutureWarning:


Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```

[30]: # replacing the forecasted price for the last 10 days
forecasted_price = pd.DataFrame(data=np.concatenate([stock_price[:100].values,
↪next_means]),
                                index=stock_price.index)

plt.figure(figsize=(12,6))
plt.plot(stock_price[len(stock_price)-100:], 'b',lw=1.5)
plt.plot(forecasted_price[len(stock_price)-100:], 'r',lw=1.5)
plt.legend(['Stock Price', 'Kalman Filter Forecasted Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price for 100
↪Days');

```



```
[31]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
      ↪mean_squared_error, mean_absolute_error
import math

_y_stock = stock_price[len(stock_price)-100:] # changed from 10 to 100, after
      ↪experiment put it back to 10
kalman_forecast_y = forecasted_price[len(stock_price)-100:] # changed from 10
      ↪to 100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
print("Mean absolute percentage error {}".
      ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
      ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
print("Mean absolute error {}".format(mean_absolute_error(_y_stock,
      ↪kalman_forecast_y)))
```

R square 0.9041106970476439

Mean absolute percentage error 0.031062199149145875

Root Mean Square Error RMSE 134.7597223756434

Mean absolute error 95.2399569958813

```
[32]: kalmanFilter = KalmanFilter(transition_matrices = [1],
      observation_matrices = [1],
      initial_state_mean = 0,
      initial_state_covariance = 1,
      observation_covariance=1,
      transition_covariance=.0001)

# Predicting for ETH-USD
stock_price = data['Adj Close']['ETH-USD']

# Kalman Forecast for Next 100 days
mean, cov = kalmanFilter.em(stock_price[:-50], n_iter=10).smooth(stock_price[:
      ↪-50]) # changed from 10 to 100, after experiment put it back to 10
# mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

next_means = []
next_covs = []
next_mean = mean[-1]
next_cov = cov[-1]
for i in range(50):
    next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
      ↪stock_price[-(50 - i * 1)]) # changed from 10 to 100, after experiment put
      ↪it back to 10
    next_means.append(next_mean[0])
    next_covs.append(next_cov[0])
```

```


# replacing the forecasted price for the last 10 days
forecasted_price = pd.DataFrame(data=np.concatenate([stock_price[:len(stock_price)-50].values,
↪next_means]),
                                index=stock_price.index)

plt.figure(figsize=(12,6))
plt.plot(stock_price[len(stock_price)-50:], 'b', lw=1.5)
plt.plot(forecasted_price[len(stock_price)-50:], 'r', lw=1.5)
plt.legend(['Stock Price', 'Kalman Filter Forecasted Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price for 50_
↪days');

```

<ipython-input-32-70c97d46fae8>:21: FutureWarning:

Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`



```

[33]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↪mean_squared_error, mean_absolute_error
import math

```



```

_y_stock = stock_price[len(stock_price)-50:] # changed from 10 to 100, after
↳experiment put it back to 10
kalman_forecast_y = forecasted_price[len(stock_price)-50:] # changed from 10 to
↳100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
print("Mean absolute error {}".format(mean_absolute_error(_y_stock,
↳kalman_forecast_y)))

```

R square 0.8698406878436127

Mean absolute percentage error 0.033123361785285604

Root Mean Square Error RMSE 135.19253245722786

Mean absolute error 93.34666907382754

```

[86]: kalmanFilter = KalmanFilter(transition_matrices = [1],
    observation_matrices = [1],
    initial_state_mean = 0,
    initial_state_covariance = 1,
    observation_covariance=1,
    transition_covariance=.01)

# Predicting for ETH-USD
stock_price = data['Adj Close']['ETH-USD']

# Kalman Forecast for Next 100 days
mean, cov = kalmanFilter.em(stock_price[:-10], n_iter=10).smooth(stock_price[:
↳-10]) # changed from 10 to 100, after experiment put it back to 10
# mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

next_means = []
next_covs = []
next_mean = mean[-1]
next_cov = cov[-1]
for i in range(10):
    next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
↳stock_price[-(10 - i * 1)]) # changed from 10 to 100, after experiment put
↳it back to 10
    next_means.append(next_mean[0])
    next_covs.append(next_cov[0])

# replacing the forecasted price for the last 10 days

```

```


forecasted_price = pd.DataFrame(data=np.concatenate([stock_price[:-10].values,
↳next_means]),
                                index=stock_price.index)

plt.figure(figsize=(12,6))
plt.plot(stock_price[len(stock_price)-10:], 'b',lw=1.5)
plt.plot(forecasted_price[len(stock_price)-10:], 'r',lw=1.5)
plt.legend(['Stock Price', 'Kalman Filter Forecasted Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price for 10_
↳days');

```

<ipython-input-86-ebbd399fcbf2>:21: FutureWarning:

Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`



```

[87]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↳mean_squared_error, mean_absolute_error
import math

_y_stock = stock_price[len(stock_price)-10:] # changed from 10 to 100, after
↳experiment put it back to 10

```

```

kalman_forecast_y = forecasted_price[len(stock_price)-10:] # changed from 10 to
↳100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
print("Mean absolute error {}".format(mean_absolute_error(_y_stock,
↳kalman_forecast_y)))

```

R square 0.9124687521208008

Mean absolute percentage error 0.007428851406645756

Root Mean Square Error RMSE 31.72503430765914

Mean absolute error 19.307972955498645

7 Kalman Filter Forecast for Volume

In the below section we will forecast Volume for Ethereum data for 200 days

```

[88]: kalmanFilter = KalmanFilter(transition_matrices = [1],
    observation_matrices = [1],
    initial_state_mean = 0,
    initial_state_covariance = 1,
    observation_covariance=1,
    transition_covariance=2)

# Predicting for ETH-USD
volume = data['Volume']['ETH-USD']

# Kalman Forecast for Next 100 days
mean, cov = kalmanFilter.em(volume[:200], n_iter=10).smooth(volume[:200]) #
↳changed from 10 to 100, after experiment put it back to 10
# mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

next_means = []
next_covs = []
next_mean = mean[-1]
next_cov = cov[-1]
for i in range(200):
    next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
↳volume[-(200 - i * 1)]) # changed from 10 to 100, after experiment put it
↳back to 10
    next_means.append(next_mean[0])
    next_covs.append(next_cov[0])

```

```


# replacing the forecasted price for the last 10 days
forecasted_volume = pd.DataFrame(data=np.concatenate([volume[:-200].values,
↳next_means]),
                                index=volume.index)

plt.figure(figsize=(12,6))
plt.plot(volume[len(volume)-200:], 'b',lw=1.5)
plt.plot(forecasted_volume[len(volume)-200:], 'r',lw=1.5)
plt.legend(['ETH-USD Volume', 'Kalman Filter Forecasted Volume'])
plt.xlabel('Day')
plt.ylabel('Volume')
plt.title('Comparison between Daily Forecasts and Volume');

```

<ipython-input-88-b19efca3b6aa>:21: FutureWarning:

Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`



```

[89]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↳mean_squared_error, mean_absolute_error
import math

_y_volume = volume[len(volume)-200:] # changed from 10 to 100, after experiment
↳put it back to 10

```

```

kalman_forecast_y = forecasted_volume[len(volume)-200:] # changed from 10 to
↳100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_volume, kalman_forecast_y)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(_y_volume, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(_y_volume, kalman_forecast_y))))
print("Mean absolute error {}".format(mean_absolute_error(_y_volume,
↳kalman_forecast_y)))

```

R square 0.9286127908117351

Mean absolute percentage error 0.08178317315450043

Root Mean Square Error RMSE 2049202655.6268578

Mean absolute error 1374597299.8002546

```

[83]: kalmanFilter = KalmanFilter(transition_matrices = [1],
    observation_matrices = [1],
    initial_state_mean = 0,
    initial_state_covariance = 1,
    observation_covariance=1,
    transition_covariance=2)

# Predicting for ETH-USD
volume = data['Volume']['ETH-USD']

# Kalman Forecast for Next 100 days
mean, cov = kalmanFilter.em(volume[:100], n_iter=10).smooth(volume[:100]) #
↳changed from 10 to 100, after experiment put it back to 10
# mean, cov = kalmanFilter.em(stock_price[1], n_iter=10)

next_means = []
next_covs = []
next_mean = mean[-1]
next_cov = cov[-1]
for i in range(100):
    next_mean, next_cov = kalmanFilter.filter_update(next_mean, next_cov,
↳volume[-(100 - i * 1)]) # changed from 10 to 100, after experiment put it
↳back to 10
    next_means.append(next_mean[0])
    next_covs.append(next_cov[0])


# replacing the forecasted price for the last 10 days
forecasted_volume = pd.DataFrame(data=np.concatenate([volume[:100].values,
↳next_means]),
    index=volume.index)

```

```
plt.figure(figsize=(12,6))
plt.plot(volume[len(volume)-100:] , 'b', lw=1.5)
plt.plot(forecasted_volume[len(volume)-100:] , 'r', lw=1.5)
plt.legend(['ETH-USD Volume', 'Kalman Filter Forecasted Volume'])
plt.xlabel('Day')
plt.ylabel('Volume')
plt.title('Comparison between Daily Forecasts and Volume');
```

<ipython-input-83-cc2dd358d03f>:21: FutureWarning:

Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`



```
[84]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
      ↪ mean_squared_error, mean_absolute_error
      import math

      _y_volume = volume[len(volume)-100:] # changed from 10 to 100, after experiment
      ↪ put it back to 10
      kalman_forecast_y = forecasted_volume[len(volume)-100:] # changed from 10 to
      ↪ 100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_volume, kalman_forecast_y)))
      print("Mean absolute percentage error {}".
      ↪ format(mean_absolute_percentage_error(_y_volume, kalman_forecast_y)))
```

```
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(_y_volume, kalman_forecast_y))))
print("Mean absolute error {}".format(mean_absolute_error(_y_volume,
↳kalman_forecast_y)))
```

```
R square 0.9436112773958336
Mean absolute percentage error 0.07167967220019686
Root Mean Square Error RMSE 1813420368.38731
Mean absolute error 1144993804.7182558
```

8 Using Traditional ML for Prediction

In this section we will explore how traditional methods work for stock price prediction. We will be using multivariable regression technique to predict closing price of the stock for the next day. The feature engineering for this model will include column values high, low, open, close, volume to determine the Adjusted close price for the next day.

```
[40]: data.head()
```

```
[40]: Price                               Adj Close                               Close \
Ticker                               BTC-USD                               ETH-USD                               BTC-USD
Date
2022-09-01 00:00:00+00:00  20127.140625  1586.176758  20127.140625
2022-09-02 00:00:00+00:00  19969.771484  1577.220459  19969.771484
2022-09-03 00:00:00+00:00  19832.087891  1556.872681  19832.087891
2022-09-04 00:00:00+00:00  19986.712891  1577.641602  19986.712891
2022-09-05 00:00:00+00:00  19812.371094  1617.183228  19812.371094
```

```
Price                               High \
Ticker                               ETH-USD                               BTC-USD                               ETH-USD
Date
2022-09-01 00:00:00+00:00  1586.176758  20198.390625  1593.082764
2022-09-02 00:00:00+00:00  1577.220459  20401.568359  1643.183228
2022-09-03 00:00:00+00:00  1556.872681  20037.009766  1579.454346
2022-09-04 00:00:00+00:00  1577.641602  19999.689453  1578.009277
2022-09-05 00:00:00+00:00  1617.183228  20031.160156  1621.661377
```

```
Price                               Low                               Open \
Ticker                               BTC-USD                               ETH-USD                               BTC-USD
Date
2022-09-01 00:00:00+00:00  19653.968750  1520.188354  20050.498047
2022-09-02 00:00:00+00:00  19814.765625  1551.877930  20126.072266
2022-09-03 00:00:00+00:00  19698.355469  1541.672119  19969.718750
2022-09-04 00:00:00+00:00  19636.816406  1543.698853  19832.470703
2022-09-05 00:00:00+00:00  19673.046875  1559.781860  19988.789062
```

```
Price                               Volume
```

Ticker		ETH-USD	BTC-USD	ETH-USD
Date				
2022-09-01 00:00:00+00:00		1553.756348	30182031010	16434276817
2022-09-02 00:00:00+00:00		1586.017944	29123998928	17708478709
2022-09-03 00:00:00+00:00		1577.213745	23613051457	9516825994
2022-09-04 00:00:00+00:00		1556.895874	25245861652	8884144998
2022-09-05 00:00:00+00:00		1577.884033	28813460025	13060541168

```
[41]: btc_data = {'Open': data['Open']['BTC-USD'],
                  'Close': data['Close']['BTC-USD'],
                  'High': data['High']['BTC-USD'],
                  'Low': data['Low']['BTC-USD'],
                  'Volume': data['Low']['BTC-USD'],
                  'Adj Close': data['Adj Close']['BTC-USD']}

eth_data = {'Open': data['Open']['ETH-USD'],
            'Close': data['Close']['ETH-USD'],
            'High': data['High']['ETH-USD'],
            'Low': data['Low']['ETH-USD'],
            'Volume': data['Low']['ETH-USD'],
            'Adj Close': data['Adj Close']['ETH-USD']}

btc_df = pd.DataFrame(btc_data, columns = ['Open', 'Close', 'High', 'Low', 'Volume', 'Adj Close'])
eth_df = pd.DataFrame(eth_data, columns = ['Open', 'Close', 'High', 'Low', 'Volume', 'Adj Close'])
```

```
[42]: eth_df.head()
```

```
[42]:
```

		Open	Close	High	Low \
Date					
2022-09-01 00:00:00+00:00		1553.756348	1586.176758	1593.082764	1520.188354
2022-09-02 00:00:00+00:00		1586.017944	1577.220459	1643.183228	1551.877930
2022-09-03 00:00:00+00:00		1577.213745	1556.872681	1579.454346	1541.672119
2022-09-04 00:00:00+00:00		1556.895874	1577.641602	1578.009277	1543.698853
2022-09-05 00:00:00+00:00		1577.884033	1617.183228	1621.661377	1559.781860

		Volume	Adj Close
Date			
2022-09-01 00:00:00+00:00		1520.188354	1586.176758
2022-09-02 00:00:00+00:00		1551.877930	1577.220459
2022-09-03 00:00:00+00:00		1541.672119	1556.872681
2022-09-04 00:00:00+00:00		1543.698853	1577.641602
2022-09-05 00:00:00+00:00		1559.781860	1617.183228


```
[43]: eth_df.describe()
```

```
[43]:
```

	Open	Close	High	Low	Volume \
count	730.000000	730.000000	730.000000	730.000000	730.000000
mean	2153.339377	2154.649833	2197.880829	2107.257615	2107.257615
std	761.057358	760.858790	781.196683	736.665024	736.665024
min	1100.107178	1100.169800	1136.442627	1081.138184	1081.138184
25%	1617.854645	1620.496521	1644.864014	1580.678528	1580.678528
50%	1865.844604	1866.100159	1887.949524	1845.784241	1845.784241
75%	2640.590576	2640.965576	2709.429199	2585.498413	2585.498413
max	4066.690430	4066.445068	4092.284180	3936.627197	3936.627197

	Adj Close
count	730.000000
mean	2154.649833
std	760.858790
min	1100.169800
25%	1620.496521
50%	1866.100159
75%	2640.965576
max	4066.445068

```
[44]: btc_df.describe()
```

```
[44]:
```

	Open	Close	High	Low	Volume \
count	730.000000	730.000000	730.000000	730.000000	730.000000
mean	37464.146091	37518.319863	38162.137949	36780.942348	36780.942348
std	17678.426052	17685.270401	18083.910703	17220.739292	17220.739292
min	15782.300781	15787.284180	16253.047852	15599.046875	15599.046875
25%	23627.717285	23665.855469	24119.581543	23253.754883	23253.754883
50%	29403.917969	29412.204102	29845.836914	29113.966797	29113.966797
75%	55644.687500	55988.014648	57679.622070	54234.083008	54234.083008
max	73079.375000	73083.500000	73750.070312	71334.093750	71334.093750

	Adj Close
count	730.000000
mean	37518.319863
std	17685.270401
min	15787.284180
25%	23665.855469
50%	29412.204102
75%	55988.014648
max	73083.500000

```
[45]: btc_df.head()
```

```
[45]:
```

Date	Open	Close	High \
2022-09-01 00:00:00+00:00	20050.498047	20127.140625	20198.390625
2022-09-02 00:00:00+00:00	20126.072266	19969.771484	20401.568359
2022-09-03 00:00:00+00:00	19969.718750	19832.087891	20037.009766
2022-09-04 00:00:00+00:00	19832.470703	19986.712891	19999.689453
2022-09-05 00:00:00+00:00	19988.789062	19812.371094	20031.160156

Date	Low	Volume	Adj Close
2022-09-01 00:00:00+00:00	19653.968750	19653.968750	20127.140625
2022-09-02 00:00:00+00:00	19814.765625	19814.765625	19969.771484
2022-09-03 00:00:00+00:00	19698.355469	19698.355469	19832.087891
2022-09-04 00:00:00+00:00	19636.816406	19636.816406	19986.712891
2022-09-05 00:00:00+00:00	19673.046875	19673.046875	19812.371094

```
[46]: eth_df['Target'] = eth_df['Adj Close'].shift(-1)
eth_df.dropna(inplace=True)
eth_df.head()
```

```
[46]:
```

Date	Open	Close	High	Low \
2022-09-01 00:00:00+00:00	1553.756348	1586.176758	1593.082764	1520.188354
2022-09-02 00:00:00+00:00	1586.017944	1577.220459	1643.183228	1551.877930
2022-09-03 00:00:00+00:00	1577.213745	1556.872681	1579.454346	1541.672119
2022-09-04 00:00:00+00:00	1556.895874	1577.641602	1578.009277	1543.698853
2022-09-05 00:00:00+00:00	1577.884033	1617.183228	1621.661377	1559.781860

Date	Volume	Adj Close	Target
2022-09-01 00:00:00+00:00	1520.188354	1586.176758	1577.220459
2022-09-02 00:00:00+00:00	1551.877930	1577.220459	1556.872681
2022-09-03 00:00:00+00:00	1541.672119	1556.872681	1577.641602
2022-09-04 00:00:00+00:00	1543.698853	1577.641602	1617.183228
2022-09-05 00:00:00+00:00	1559.781860	1617.183228	1561.748535

```
[47]: eth_df.tail()
```

```
[47]:
```

Date	Open	Close	High	Low \
2024-08-25 00:00:00+00:00	2769.098145	2749.157715	2793.012939	2736.088867
2024-08-26 00:00:00+00:00	2749.247559	2681.340576	2763.004150	2668.886719
2024-08-27 00:00:00+00:00	2681.622803	2458.726562	2700.152832	2401.175049
2024-08-28 00:00:00+00:00	2458.904785	2528.415527	2553.820068	2422.293701
2024-08-29 00:00:00+00:00	2528.362305	2528.792725	2595.977051	2507.502441

Volume	Adj Close	Target
--------	-----------	--------

Date				
2024-08-25 00:00:00+00:00	2736.088867	2749.157715	2681.340576	
2024-08-26 00:00:00+00:00	2668.886719	2681.340576	2458.726562	
2024-08-27 00:00:00+00:00	2401.175049	2458.726562	2528.415527	
2024-08-28 00:00:00+00:00	2422.293701	2528.415527	2528.792725	
2024-08-29 00:00:00+00:00	2507.502441	2528.792725	2525.822021	

We can define as the forecast set for the last 10 data points to compare with our Kalman Forecast.

```
[48]: forecast_set = eth_df[-200:] # validation set
      training_set = eth_df[:-200] # changed from 10 to 100, after experiment put it
      ↪ back to 10
```

```
[49]: forecast_set
```

```
[49]:
```

	Open	Close	High	Low \
Date				
2024-02-12 00:00:00+00:00	2507.578857	2658.115967	2663.842773	2473.812012
2024-02-13 00:00:00+00:00	2659.586182	2642.185303	2686.455078	2599.169434
2024-02-14 00:00:00+00:00	2641.685303	2777.902344	2786.893555	2621.025391
2024-02-15 00:00:00+00:00	2777.601318	2824.378906	2865.845459	2764.010498
2024-02-16 00:00:00+00:00	2825.480713	2803.691406	2858.450439	2760.331055
...
2024-08-25 00:00:00+00:00	2769.098145	2749.157715	2793.012939	2736.088867
2024-08-26 00:00:00+00:00	2749.247559	2681.340576	2763.004150	2668.886719
2024-08-27 00:00:00+00:00	2681.622803	2458.726562	2700.152832	2401.175049
2024-08-28 00:00:00+00:00	2458.904785	2528.415527	2553.820068	2422.293701
2024-08-29 00:00:00+00:00	2528.362305	2528.792725	2595.977051	2507.502441

	Volume	Adj Close	Target
Date			
2024-02-12 00:00:00+00:00	2473.812012	2658.115967	2642.185303
2024-02-13 00:00:00+00:00	2599.169434	2642.185303	2777.902344
2024-02-14 00:00:00+00:00	2621.025391	2777.902344	2824.378906
2024-02-15 00:00:00+00:00	2764.010498	2824.378906	2803.691406
2024-02-16 00:00:00+00:00	2760.331055	2803.691406	2786.672607
...
2024-08-25 00:00:00+00:00	2736.088867	2749.157715	2681.340576
2024-08-26 00:00:00+00:00	2668.886719	2681.340576	2458.726562
2024-08-27 00:00:00+00:00	2401.175049	2458.726562	2528.415527
2024-08-28 00:00:00+00:00	2422.293701	2528.415527	2528.792725
2024-08-29 00:00:00+00:00	2507.502441	2528.792725	2525.822021

[200 rows x 7 columns]

```
[50]: # len(X_test)
```

```
[51]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing

X = training_set.drop('Target', axis=1)
y = training_set['Target']

# calling pre-processing for scaling or standardization of the X values
X_scaled = preprocessing.scale(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
↳random_state=42, shuffle=False)

lr = LinearRegression(n_jobs=-1)
lr.fit(X_train, y_train)

score = lr.score(X_test, y_test)
print('The linear regression confidence is {}'.format(score))
```

The linear regression confidence is 0.8901272092118377

```
[52]: prediction = lr.predict(X_test[:100])

# from sklearn.metrics import r2_score , mean_absolute_percentage_error
# print("R square {}".format(r2_score(y_test, prediction)))
# print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(y_test, prediction)))
```

```
[53]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↳mean_squared_error, mean_absolute_error
import math

print("R square {}".format(r2_score(y_test[:100], prediction)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(y_test[:100], prediction)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(y_test[:100], prediction))))
print("Mean absolute error {}".format(mean_absolute_error(y_test[:100],
↳prediction)))
```


R square 0.8852723302985438

Mean absolute percentage error 0.020247466438748578

Root Mean Square Error RMSE 64.68978629963057

Mean absolute error 45.34421937020249

```
[54]: plt.figure(figsize=(12,6))
plt.plot(y_test[:100] , 'b', lw=1.5)
plt.plot(pd.DataFrame(data=prediction, index=y_test[:100].index) , 'r', lw=1.5)
plt.legend(['Actual Stock Price', 'Linear Regression - Predicted Stock Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price - Test_
↳Set');
```



```
[55]: prediction = lr.predict(X_test[:50])
from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↳mean_squared_error, mean_absolute_error
import math

print("R square {}".format(r2_score(y_test[:50], prediction)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(y_test[:50], prediction)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(y_test[:50], prediction))))
print("Mean absolute error {}".format(mean_absolute_error(y_test[:50],
↳prediction)))
```


R square 0.8508220091574861

Mean absolute percentage error 0.02041025364721451

Root Mean Square Error RMSE 60.52601272456529

Mean absolute error 42.913583621540994

```
[56]: plt.figure(figsize=(12,6))
plt.plot(y_test[:50] , 'b', lw=1.5)
plt.plot(pd.DataFrame(data=prediction, index=y_test[:50].index) , 'r', lw=1.5)
plt.legend(['Actual Stock Price', 'Linear Regression - Predicted Stock Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price - Test_
↳Set');
```



```
[57]: prediction = lr.predict(X_test[:10])
from sklearn.metrics import r2_score , mean_absolute_percentage_error,
↳mean_squared_error, mean_absolute_error
import math

print("R square {}".format(r2_score(y_test[:10], prediction)))
print("Mean absolute percentage error {}".
↳format(mean_absolute_percentage_error(y_test[:10], prediction)))
print("Root Mean Square Error RMSE {}".format(math.
↳sqrt(mean_squared_error(y_test[:10], prediction))))
print("Mean absolute error {}".format(mean_absolute_error(y_test[:10],
↳prediction)))
```


R square 0.4464668178124982

Mean absolute percentage error 0.012516394619436188

Root Mean Square Error RMSE 26.64129516909742

Mean absolute error 23.124079376954207

```
[58]: plt.figure(figsize=(12,6))
plt.plot(y_test[:10] , 'b', lw=1.5)
plt.plot(pd.DataFrame(data=prediction, index=y_test[:10].index) , 'r', lw=1.5)
plt.legend(['Actual Stock Price', 'Linear Regression - Predicted Stock Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price - Test_
↵Set');
```



[58]:

```
[59]: # We will not forecast for the next days based on this model
forecast_X = forecast_set.drop('Target', axis=1)
forecast_y = forecast_set['Target']
forecast_pred = lr.predict(preprocessing.scale(forecast_X))
forecast_pred
```

```
[59]: array([1217.91571333, 1216.1973982 , 1322.02023933, 1370.38365189,
1354.01390365, 1335.68814119, 1411.43525403, 1469.86042295,
1523.72327363, 1490.37171101, 1495.12993538, 1457.19613254,
1509.49103753, 1608.42146148, 1663.69507854, 1726.4507585 ,
1842.10495806, 1819.31722747, 1885.72513323, 1880.2195351 ,
1929.57281682, 2044.33892688, 1981.11571383, 2200.08262905,
2257.88224042, 2281.10149117, 2298.90060741, 2268.37202312,
2404.52098619, 2347.55004232, 2375.99336318, 2264.35578722,
2142.2916627 , 1974.17046562, 2051.91177114, 1962.26772122,
1673.80850231, 1926.04722158, 1939.10137169, 1809.04417871,
```

```

1808.22907973, 1897.52621746, 2015.49548376, 2022.81446159,
1952.10872937, 1994.23539975, 1957.02888573, 1955.11849923,
2062.61473059, 1950.07988573, 1762.85193517, 1781.18522494,
1802.83978859, 1784.97151683, 1821.82626384, 1899.48239066,
2093.80131211, 1957.44444349, 1974.22140874, 1954.81716688,
1731.69851211, 1528.52931917, 1637.63449033, 1611.71757956,
1588.55291317, 1510.06876242, 1571.25371138, 1559.77248928,
1646.14150478, 1647.21106359, 1689.83449389, 1705.81581857,
1645.82738758, 1649.8785952 , 1631.55327097, 1723.95362221,
1749.32809718, 1699.72533838, 1535.78015099, 1485.16888968,
1505.33322301, 1600.49484344, 1622.87813586, 1635.15264144,
1582.33607956, 1533.60429276, 1499.5310146 , 1547.14179316,
1449.4086738 , 1446.2221755 , 1460.34258199, 1474.4021061 ,
1422.81302033, 1540.95390057, 1477.77928282, 1591.77010462,
1624.8116187 , 1584.49700168, 2044.43499103, 2183.46310097,
2144.73023246, 2173.80887091, 2136.5061315 , 2156.80553927,
2219.84371267, 2279.66788728, 2234.10786577, 2174.40223398,
2156.87148163, 2169.70571118, 2208.87095159, 2185.45848431,
2177.01190551, 2207.41429026, 2250.80190822, 2210.21490447,
2100.78121042, 2100.79584598, 2119.2508682 , 2088.57290852,
1947.74702974, 1995.68538769, 1921.2689725 , 1923.06757158,
1997.02770605, 2044.584827 , 1958.34944967, 1924.36412948,
1991.52915793, 1959.94648173, 1955.90051719, 1941.5636446 ,
1881.46883844, 1814.04420361, 1854.42590641, 1834.76857703,
1895.19029588, 1843.85905331, 1839.82879103, 1884.60458989,
1899.06525869, 1876.09982191, 1773.59720704, 1580.2533993 ,
1499.26460873, 1572.85955284, 1467.89462867, 1526.94796516,
1574.18623187, 1605.31957206, 1609.89494555, 1629.92477652,
1668.78512425, 1724.66504 , 1919.93513024, 1893.26514808,
1858.12843421, 1883.82554778, 1943.73890354, 1961.0376562 ,
1967.95165717, 1900.50979209, 1927.9188702 , 1811.80603737,
1670.43224362, 1749.31916994, 1731.98801923, 1747.23478852,
1794.05344506, 1759.3973609 , 1723.31500299, 1684.28041568,
1518.97829078, 1441.53056507, 1263.16346476, 1014.66780036,
1064.54911545, 967.379339 , 1229.57357063, 1182.84953281,
1189.74094249, 1150.66861676, 1273.53733526, 1263.77255521,
1239.24177516, 1156.93062278, 1174.93204633, 1192.73181538,
1195.90344294, 1208.21075628, 1164.05984206, 1203.18748243,
1199.17668029, 1313.6086976 , 1326.00771347, 1309.57842541,
1253.6660973 , 1068.45111552, 1114.32398562, 1123.44135779])

```

```
[60]: forecast_y
```

```
[60]: Date
```

```

2024-02-12 00:00:00+00:00    2642.185303
2024-02-13 00:00:00+00:00    2777.902344
2024-02-14 00:00:00+00:00    2824.378906

```



```


2024-02-15 00:00:00+00:00    2803.691406
2024-02-16 00:00:00+00:00    2786.672607
...
2024-08-25 00:00:00+00:00    2681.340576
2024-08-26 00:00:00+00:00    2458.726562
2024-08-27 00:00:00+00:00    2528.415527
2024-08-28 00:00:00+00:00    2528.792725
2024-08-29 00:00:00+00:00    2525.822021
Name: Target, Length: 200, dtype: float64

```

```

[61]: plt.figure(figsize=(12,6))
plt.plot(forecast_y , 'b',lw=1.5)
plt.plot(pd.DataFrame(data=forecast_pred, index=forecast_y.index) , 'r',lw=1.5)
plt.legend(['Actual Stock Price', 'Linear Regression - Forecasted Stock Price'])
plt.xlabel('Day')
plt.ylabel('Stock Price')
plt.title('Comparison between Daily Forecasts and Original Stock Price - Last 10 data points');

```



Metrics from LR for the stock price

```

[62]: from sklearn.metrics import r2_score , mean_absolute_percentage_error,
      ↪ mean_squared_error, mean_absolute_error
import math

print("R square {}".format(r2_score(forecast_y, forecast_pred)))

```

```

print("Mean absolute percentage error {}".
      ↪format(mean_absolute_percentage_error(forecast_y, forecast_pred)))
print("Root Mean Square Error RMSE {}".format(math.
      ↪sqrt(mean_squared_error(forecast_y, forecast_pred))))
print("Mean absolute error {}".format(mean_absolute_error(forecast_y,
      ↪forecast_pred)))

```

R square -14.678373150650259
 Mean absolute percentage error 0.470722951756699
 Root Mean Square Error RMSE 1523.4539014820086
 Mean absolute error 1518.8060010609427

Metrics from the Kalman Filter forecast

```

[63]: _y_stock = stock_price[len(stock_price)-2000:] # changed from 10 to 100, after
      ↪experiment put it back to 10
kalman_forecast_y = forecasted_price[len(stock_price)-2000:] # changed from 10
      ↪to 100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
print("Mean absolute percentage error MAPE {}".
      ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
      ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
      ↪kalman_forecast_y)))

```

R square 0.9998118018583926
 Mean absolute percentage error MAPE 0.00037247146220583305
 Root Mean Square Error RMSE 10.430714570659806
 Mean absolute error MAE 0.966728266157691

```

[64]: _y_stock = stock_price[len(stock_price)-500:] # changed from 10 to 100, after
      ↪experiment put it back to 10
kalman_forecast_y = forecasted_price[len(stock_price)-500:] # changed from 10
      ↪to 100, after experiment put it back to 10
print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
print("Mean absolute percentage error MAPE {}".
      ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
print("Root Mean Square Error RMSE {}".format(math.
      ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
      ↪kalman_forecast_y)))

```

R square 0.9996946969956958
 Mean absolute percentage error MAPE 0.0005438083348205162
 Root Mean Square Error RMSE 12.60348036947252
 Mean absolute error MAE 1.4114232685902288

```
[65]: _y_stock = stock_price[len(stock_price)-200:] # changed from 10 to 100, after
      ↪experiment put it back to 10
      kalman_forecast_y = forecasted_price[len(stock_price)-200:] # changed from 10
      ↪to 100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
      print("Mean absolute percentage error MAPE {}".
            ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
      print("Root Mean Square Error RMSE {}".format(math.
            ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
      print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
            ↪kalman_forecast_y)))
```

R square 0.9973173541119791
Mean absolute percentage error MAPE 0.0013595208370512906
Root Mean Square Error RMSE 19.92785220637683
Mean absolute error MAE 3.5285581714755723

```
[66]: _y_stock = stock_price[len(stock_price)-100:] # changed from 10 to 100, after
      ↪experiment put it back to 10
      kalman_forecast_y = forecasted_price[len(stock_price)-100:] # changed from 10
      ↪to 100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
      print("Mean absolute percentage error MAPE {}".
            ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
      print("Root Mean Square Error RMSE {}".format(math.
            ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
      print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
            ↪kalman_forecast_y)))
```

R square 0.9958062655293043
Mean absolute percentage error MAPE 0.0027190416741025808
Root Mean Square Error RMSE 28.182238859224718
Mean absolute error MAE 7.0571163429511445

```
[67]: _y_stock = stock_price[len(stock_price)-50:] # changed from 10 to 100, after
      ↪experiment put it back to 10
      kalman_forecast_y = forecasted_price[len(stock_price)-50:] # changed from 10 to
      ↪100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
      print("Mean absolute percentage error MAPE {}".
            ↪format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
      print("Root Mean Square Error RMSE {}".format(math.
            ↪sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
      print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
            ↪kalman_forecast_y)))
```

R square 0.9886877025379025
Mean absolute percentage error MAPE 0.005438083348205162

Root Mean Square Error RMSE 39.85570441275366
Mean absolute error MAE 14.114232685902289

```
[68]: _y_stock = stock_price[len(stock_price)-10:] # changed from 10 to 100, after
      ↪ experiment put it back to 10
      kalman_forecast_y = forecasted_price[len(stock_price)-10:] # changed from 10 to
      ↪ 100, after experiment put it back to 10
      print("R square {}".format(r2_score(_y_stock, kalman_forecast_y)))
      print("Mean absolute percentage error MAPE {}".
      ↪ format(mean_absolute_percentage_error(_y_stock, kalman_forecast_y)))
      print("Root Mean Square Error RMSE {}".format(math.
      ↪ sqrt(mean_squared_error(_y_stock, kalman_forecast_y))))
      print("Mean absolute error MAE {}".format(mean_absolute_error(_y_stock,
      ↪ kalman_forecast_y)))
```

R square 0.30926748192137554
Mean absolute percentage error MAPE 0.027190416741025812
Root Mean Square Error RMSE 89.12006435805553
Mean absolute error MAE 70.57116342951144